

Neural Network design

In addition to learning *how* to use Neural Networks, we hope this course has stimulated your curiosity.

What strikes me as particularly curious:


- Neural Networks create representations that make it easier for the Head Layer to solve a Classical ML task
- We don't know *how* these representations are created
- We certainly have not given an explicit instruction or requirement
- We may not even *know* how to interpret these representation
- Yet the representation seems to be both useful for a particular task
- And *Transferable* to other tasks

Just as a Neural Network for a vision task


- Seems to learn concepts
 - "Dimensions of meaning"

Features by layer



Layer 1



Layer 2



Layer 3



So too do Word Embeddings seem to learn dimensions of meaning.

Meaning was learned without direction: pretty remarkable !

Another curiosity: we are able to solve complex tasks (e.g. NLP, image recognition) with *simple* programs.

The "art" of Neural Networks is *not* highly skilled programming but instead

- Being clever and diligent in acquiring enough training examples
- Creating a Loss Function that captures the essence of the problem

For example: Neural Style Transfer is a task that

- Takes one image (the "Content Image")
- And an artistic style, as expressed by a "Style Image"
- Produces a new image that re-expresses the Content Image in the style of the Style Image

Content Image



Style Image



Generated Image



The "trick" in solving this task is in writing the Loss function

- Not in designing the network
- With the Loss function in hand, we then apply the skills we learned to minimize Loss Functions
- And the task is solved

Without going into detail the Loss Function has two parts

- A "content loss": the generated image $\vec{\mathbf{x}}$ should be close to the Source Image $\vec{\mathbf{p}}$
- A "style loss": the style of the generated image $\vec{\mathbf{x}}$ and the Style Image $\vec{\mathbf{a}}$ should be close

$$\mathcal{L} = \mathcal{L}_{\text{content}}(\vec{\mathbf{p}}, \vec{\mathbf{x}}) + \mathcal{L}_{\text{style}}(\vec{\mathbf{a}}, \vec{\mathbf{x}})$$

To be sure: there is some cleverness involved in

- Defining what "style" is
- What is the best measure of "being close"

but given the framework of the Loss Function, these tasks are closer to Engineering than Art.

When these skills are combined with Transfer Learning

- You are truly able to "stand on the shoulder of giants"
- And hopefully solve those tasks that are meaningful to your domain

```
In [ ]: print("Done")
```