

## Context Sensitive Memory

A Context Sensitive Memory is like a Python dict:

- data is stored via (key, value) pairs
- a "query" matches a key and returns the associated value

The difference from a Python dict

- the query is compared to every key, and a "weight" indicating strength of match is returned
  - match can be approximate
- the value returned is the weighted sum of all values
  - if there is an exact match of one and only one key, this is equivalent to a Python dict.

## Context Sensitive Memory

- a collect of key/value pairs, like a Python dict

$$M = \{(k_t, v_t) | 1 \leq t \leq T\}$$

- lookup: pass in a "query", get a value-like output

As we learned in studying gates: the lookup needs to make soft choices rather than hard choices to be differentiable.

## Normalized scores

$$\alpha(q, k) = \frac{\exp(\text{score}(q, k))}{\sum_{k' \in \text{keys}(M)} \exp(\text{score}(q, k'))}$$

## Soft lookup

$$\mathbf{c} = \text{lookup}(q, M) = \sum_{(k,v) \in M} \alpha(q, k) * v$$

## Scoring functions

Redefine using generic k,v rather than h\_t

There are several choices for the scoring function

$$\text{score}(\mathbf{h}_{(t)}, \bar{\mathbf{h}}_{(t')}) = \begin{cases} \mathbf{h}_{(t)}^T \cdot \bar{\mathbf{h}}_{(t')} & \text{dot product, cosine similarity} \\ \mathbf{h}_{(t)}^T \mathbf{W}_\alpha \bar{\mathbf{h}}_{(t')} & \text{general} \\ \mathbf{v}_\alpha^T \tanh(\mathbf{W}_\alpha [\mathbf{h}_{(t)}; \bar{\mathbf{h}}_{(t')}]) & \text{concat} \end{cases}$$

**Note**

What is  $\mathbf{v}_\alpha^T$  ?

```


decode_init(enc_states): h = 0 self.s = enc_states return h, s
def decode_step(h, x, s): """ h: hidden state (t-1)
x: input t teacher forcing: x == y_t inference : x =  $\hat{y}_{(t-1)}$ 
s: array of encoder hidden states """ #
Update hidden state, based on input x h, out = RNN(h, x) # Compute attention weights # query == h (new state of decoder)
# key == value == s att_weights = ATT(h, s) # Compute context c = att_weights * s #
 $\hat{y}$  is function of h and c (rather than just h as in NN w/o attention) y = g( [h,c] ) return h, y
def decode(enc_states, y=None ): h, s = decode_init(enc_states) for t in range(1, t): h, y_hat = decode_step(h, x, s) #
Create next input as output of this time step if y is not None: # Training: teacher forces output to be correct answer
x = y[t] else: # Test: output (which becomes next input) x = y_hat

```

# Attention

Consider a many to many implementation of a Recurrent NN (RNN, LSTM, etc).

## RNN Encoder/Decoder



An example might be a network that adds descriptions/captions to a stream of images (video)

- input sequence: a sequence of frames
- output sequence: a sequence of words

or that translates from one language to another

- input sequence: words in source language
- output sequence: words in target language

It is very possible that the next word (time step  $t$ ) might refer to a much earlier frame ( $t' < t$ ).

A similar thing happens when translating between languages.

There is not necessarily a correspondence between output  $t$  and input  $t$ .




So an LSTM needs to decide which part of the past to "attend" (pay attention) to.

We can help it via a mechanism know as "attention", which we sketch below.

Sequence to Sequence: training (teacher forcing) + inference: No attention

Encoder



$\mathbf{x}_{(1)} \mathbf{x}_{(2)} \dots \mathbf{x}_{(\bar{T})}$


$$\begin{aligned} \mathbf{h}_{(0)} &= \bar{\mathbf{h}}_{(T)} \\ \mathbf{s} &= \bar{\mathbf{h}}_{(\bar{T})} \end{aligned}$$

Decoder

$\hat{\mathbf{y}}_{(1)} \hat{\mathbf{y}}_{(2)} \dots \hat{\mathbf{y}}_{(t)} \dots \hat{\mathbf{y}}_{(\bar{T})}$



$\mathbf{h}_{(t)}, \mathbf{s}$



$\langle \text{Start} \rangle \mathbf{y}_{(1)} \mathbf{y}_{(2)} \dots \mathbf{y}_{\bar{T}}$


Inference

$\langle \text{Start} \rangle \hat{\mathbf{y}}_{(1)} \hat{\mathbf{y}}_{(2)} \dots \hat{\mathbf{y}}_{\bar{T}}$

Training


**Decoder**

$$\hat{\mathbf{y}}_{(1)}, \hat{\mathbf{y}}_{(2)}, \dots, \hat{\mathbf{y}}_{(\bar{T})}$$




$$\mathbf{h}_{(t)}, \bar{\mathbf{h}}_{(T)}$$

$$\bar{\mathbf{h}}_{(1)}, \bar{\mathbf{h}}_{(2)}, \dots, \bar{\mathbf{h}}_{(\bar{T})}$$

**Decoder**

Sequence to Sequence: attention



The decoder is able to "select one" of the prior states, rather than just the latest one.

Of course, by now, we understand that this is a "soft" select (case/switch)

- needs to be differentiable
- so it provides a weighted combination of all prior states
  - a mask that is almost OHE becomes a true "choose one"

How does the LSTM decide which of the past states to attend to ?

Same way as all Machine Learning:

- it is controlled by weights
- that are learned by training !

So Deep Learning layers are almost becoming little computers that learn their own programs !

In [ ]:



In [3]: `print("Done")`

Done