

```
In [2]: # My standard magic ! You will see this in almost all my notebooks.

from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

# Reload all modules imported with %aimport
%load_ext autoreload
%autoreload 1

%matplotlib inline
```

```
In [3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn import datasets, linear_model
from sklearn.metrics import mean_squared_error, r2_score

import recipe_helper
%aimport recipe_helper
```

## Linear models and matrix notation

Our visualization of the data suggests that a reasonable first hypothesis is a linear relation.

$$h_{\Theta}(\mathbf{x}) = \Theta_0 + \Theta_1 \mathbf{x}$$

In our toy example there is only a single feature so  $n = \|\mathbf{x}\| = 1$

so that our linear hypothesis can be re-written as

$$h_{\Theta}(\mathbf{x}') = \hat{\mathbf{y}} = \Theta^T \cdot \mathbf{x}'$$

Notice that the Cost function (optimization objective, evaluated over the training set) mirrors our Performance Measure (evaluted over the *test* set).

Hopefully, any  $\Theta$  which does a good job on the training set also does a good job on the test set.

The critical assumptions are

- that the training and test sets are samples from the same underlying distribution
- $m$  is sufficiently large so that a  $\Theta$  estimated over the training set generalizes to unseen data

For now, the Cost function and the Performance Measure happen to be identical in form

- *but* one is evaluted over the training set; the other over the test set.

## Linear Model with higher order features

Our error analysis of the toy problem suggested that a straight line was perhaps not the best fit

- positive errors in the extremes
- negative errors in the center


Perhaps a "curve" would be a better hypothesis ? What if our data is not linear ?

Here's what the first dataset looked like


```
In [4]: (xlabel, ylabel) = ("Size", "Price Premium")

# I will give you the data via a function (so I can easily alter the data in subsequent examples)
v1, a1 = 1, .005
lin = recipe_helper.Recipe_Helper(v = v1, a = a1)
X_lin, y_lin = lin.gen_data(num=50)

v2, a2 = v1, a1*2
curv = recipe_helper.Recipe_Helper(v = v2, a = a2)
X_curve, y_curve = curv.gen_data(num=50)
_ = curv.gen_plot(X_curve, y_curve, xlabel, ylabel)
```



```
In [5]: X_orig, y_orig = X_lin, y_lin  
_ = lin.gen_plot(X_orig, y_orig, xlabel, ylabel)
```



```
In [6]: _= lin.run_regress(X_orig, y_orig)
```

Coefficients:



```
[4.93224426] [[0.96836946]]
```

R-squared (test): 0.98

Root Mean squared error (test): 6.86

R-squared (train): 0.98

Root Mean squared error (train): 6.56




We will make our point by creating a similar dataset (the "curvy" dataset) that exaggerates the curvature.

```
In [7]: _= curv.run_regress(X_curve, y_curve)
```

```
Coefficients:  
[9.86448852] [[0.93673892]]
```

```
R-squared (test): 0.91  
Root Mean squared error (test): 13.73
```

```
R-squared (train): 0.91  
Root Mean squared error (train): 13.13
```



Compared to the original, the "curvy" data set has a lot more curvature



- the  $R^2$  is still over 90%
- but the Performance Metric (RMSE) is twice as big

## Curvature in a linear model

Our (first-order) linear model was

$$y = \Theta_0 + \Theta_1 x$$

We can create a *second order* linear model by adding a feature  $x^2$ :

$$y = \Theta_0 + \Theta_1 x + \Theta_2 x^2$$

$y$  is a second order polynomial, whose plot is a curve

- but it is linear in features  $x, x^2$

In other words, we are performing feature iteration

- in this case: adding the missing feature  $x^2$

Let's modify  $\mathbf{x}^{(i)}$  from a vector of length 1:

$$\mathbf{x}^{(i)} = (\mathbf{x}_1^{(i)})$$

to a vector of length 2:

$$\mathbf{x}^{(i)} = (\mathbf{x}_1^{(i)}, \mathbf{x}_1^{(i)2})$$

by adding a squared term to the vector  $\mathbf{x}^{(i)}$ , for each  $i$ .

The modified  $\mathbf{X}'$  becomes:

$$\mathbf{X} = \begin{pmatrix} 1 & \mathbf{x}_1^{(1)} & (\mathbf{x}_1^{(1)})^2 \\ 1 & \mathbf{x}_1^{(2)} & (\mathbf{x}_1^{(2)})^2 \\ \vdots & \vdots & \\ 1 & \mathbf{x}_1^{(m)} & (\mathbf{x}_1^{(m)})^2 \end{pmatrix}$$

Note that this modified  $\mathbf{X}'$  fits perfectly within our Linear hypothesis

$$\hat{\mathbf{y}} = \mathbf{X}'\Theta$$

The requirement is that the model be linear in its *features*, **not** that the features be linear  
!

In [9]: `print("Done !")`

Done !