

# Issues with text: Sparse encoding of tokens

A token is an element of the finite set  $\mathbf{V}$

- It is thus a *categorical* variable
- Which is naturally represented by a vector of length  $||\mathbf{V}||$ 
  - One Hot Encoded (OHE)

So, in theory, we already know how to perform NLP

- Encode text as a sequence of OHE vectors
- Apply techniques (e.g., RNN) specialized to sequences

This approach may be both

- Impractical
- Sub-optimal: not taking advantage of properties inherent in text

Let us suppose that we had a function that maps a token to a OHE vector

$$\text{rep} : \text{token} \mapsto \mathbb{R}^{n \setminus \mathbf{V}}$$

Thus

$$\text{rep}(\mathbf{w}_{(t)})$$

is the OHE vector for  $\mathbf{w}_{(t)} \in \mathbf{V}$

The first issue we encounter:

- $\mathbf{V}$  is big ! A decent vocabulary is easily thousands of tokens
- So the OHE is a *long* vector ( $||\mathbf{V}||$ )

Thus, OHE may not be *practical*

It is also potentially failing to take advantage of relationships between tokens

- There is clearly a relationship between tokens  
dog, dogs  
but *no* relationship between their OHE vectors  
rep(dog), rep(dogs)

OHE vectors are *sparse*

- All zero
- Except for a single element

The sparsity both wastes space and is the cause for not capturing potential relationships between words

We will subsequently demonstrate a *dense* encoding of tokens that solves both issues.

# Issues with text: variable length sequences

Another issue with text: the sequence length is variable, not constant.

The reason this may be an issue

- Classical Machine Learning models (e.g. Logistic Regression) can only deal with fixed length inputs
- The final layer in a Neural Network is usually an implementation of a Classical model

Fortunately, we can easily identify two solutions to this issue.

Both involve reducing a variable length sequence to a fixed length encoding.

Once tokens have been encoded as a sequence of numeric vectors the solutions are

- Replace the variable number of tokens by a summary statistic (sum/average) over the tokens
- Use the final state of an RNN as an encoding of the sequence



As a notational convenience we will extend  $\text{rep}$  to sequences  $\mathbf{w}$ :

$$\text{rep}(\mathbf{w}) = [\text{rep}(\mathbf{w}_{(t)}) | 1 \leq t \leq \|\mathbf{w}\|]$$

## **Traditional methods for summarizing a variable length sequence**

The simplest way to derive a fixed length encoding of a sequence is by a summarization operation.


This is the approach that was pursued in "traditional" (pre-Neural Network) NLP.

This is what a Global Pooling layer does

- Each row is a feature over multiple spatial locations (tokens in the sequence)
- Is transformed to a single value
- Preserving the feature/channel dimension

Global Pooling  
3 features over spatial locations  
to 3 features over one location

---



Pattern: "Machine Learning"

|                  |             |         |          |          |
|------------------|-------------|---------|----------|----------|
| Machine Learning | Learning is | is easy | easy not | not hard |
|------------------|-------------|---------|----------|----------|




Pattern: "Is easy"

|                  |             |         |          |          |
|------------------|-------------|---------|----------|----------|
| Machine Learning | Learning is | is easy | easy not | not hard |
|------------------|-------------|---------|----------|----------|



Pattern: "not hard"

|                  |             |         |          |          |
|------------------|-------------|---------|----------|----------|
| Machine Learning | Learning is | is easy | easy not | not hard |
|------------------|-------------|---------|----------|----------|



Global Pooling

Machine Learning is easy not hard

Machine Learning is easy not hard

Machine Learning is easy not hard

Unfortunately, the summarization (sum/average) of a sequence

- Will lose the ordering relationship among the tokens

*"Machine Learning is easy not hard"*

*"Machine Learning is hard not easy"*

both have the same summary.

We will give a brief introduction to traditional summarization techniques.

## Bag of Words (BOW): Pooling

We define a *reduction* operation CBOW

- convert a sequence  $\mathbf{w}$  of  $||\mathbf{w}||$  elements
- each element of length  $||\text{rep}(\mathbf{w}_{(t)})||$
- to a fixed length vector of length  $||\text{rep}(\mathbf{w}_{(t)})||$

This will necessarily lose token order: this method is called *Bag of Words (BOW)*

There are many operators to achieve the reduction, which we will group under the name *pooling*

### Sum/Average

$$\text{CBOW}(\mathbf{w}) = \sum_{t=1}^{||\mathbf{w}||} \text{rep}(\mathbf{w}_{(t)})$$

Since  $\mathbf{w}_{(t)}$  is a vector, the addition operation is element-wise.

So the composite vector for the sequence is the sum of the vectors of each element in the sequence.

We can easily turn the Sum into an average by dividing by  $||\mathbf{w}||$



**Count vectorization:**

In the special case that

$$\text{rep}(\mathbf{w}_{(t)}) = \text{OHE}(\mathbf{w}_{(t)})$$

$\text{CBOW}(\mathbf{w})_j$  is equal to the number of occurrences in sequence  $\mathbf{w}$  of the  $j^{\text{th}}$  word in  $\mathbf{V}$ .

This is often called *Count Vectorization*.

## TF-IDF

Count Vectorization is simple but ignores a basic fact of language

- word "importance" is often inversely correlated with frequency in  $\frac{1}{V}$

In English:

- the words "a", "the" and "is" are extremely high frequency (so high counts in most  $w$ ).
- but are so common as to convey little meaning

On the other hand, a rare word (or sequence of words) may be very distinctive ("Machine Learning").

### *Term Frequency, Inverse Document Frequency (TF-IDF)*

- is based on the idea that a word that is *infrequent* in the wide corpus
- but is frequent in a particular document in the corpus is very meaningful in the context of the document.

So a document

- in which "Machine Learning" occurred a disproportionately high (relative to the broad corpus) number of times
- is likely to indicate that the document is dealing with the subject of Machine Learning.

**Note** A similar idea is behind many Web search algorithms (Google).

TF-IDF is similar to the Count Vectorizer, but with modified counts that are the product of

- the frequency of a word within a single document
- the inverse of the frequency of the word relative to all documents
- $v$  is a word
- $d$  is a document (collection of words) in set of documents  $D$

$\text{tf}(v, d)$  = frequency of word  $v$  in document  $d$  (Term Frequency)

$\text{df}(v)$  = number of documents that contain word  $v$

$\text{idf}(v)$  =  $\log\left(\frac{|D|}{\text{df}(v)}\right) + 1$  Inverse Document Frequency

$\text{tf-idf}(v, d)$  =  $\text{tf}(v, d) * \text{idf}(v)$


---

## **Using an RNN to obtain a fixed length encoding of a variable length sequence**

As a refresher, here is our picture of the RNN API:

# RNN many to one API

---




Although we don't know exactly what  $\mathbf{h}_{(t)}$  represents, recall that

- It is a summary of the prefix of input  $\mathbf{x}$  ending at step  $t$

So a Neural Network that

- Appends a Classical layer (Logistic Regression)
- To the output of an RNN layer
- Is an effective way of obtaining a fixed length encoding of  $\mathbf{x}$





# Conclusion

Text data is characterized by

- Tokens that are categorical variables
- Variable length sequences of tokens

We presented some "obvious" methods to deal with both issues, but they clearly have limitations.

We will move beyond the limitations in a subsequent module.

In [2]: `print("Done")`

Done