

Pandas_quick_tour

January 28, 2020

```
[1]: # My standard magic ! You will see this in almost all my notebooks.

from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

# Reload all modules imported with %aimport
%load_ext autoreload
%autoreload 1

%matplotlib inline
```

1 Pandas

[VanderPlas Chapter 3, Geron notebook](#)

```
[2]: import numpy as np
import pandas as pd
```

1.1 Series

A Series is like a NumPy ndarray but with *symbolic indexing* (like a Dictionary)

```
[3]: data = pd.Series([0.25, 0.5, 0.75, 1.0],
                      index=['a', 'b', 'c', 'd'])
data
```

```
[3]: a    0.25
     b    0.50
     c    0.75
     d    1.00
     dtype: float64
```

You select elements by using members of the index

```
[4]: print("data at 'b': ", data["b"])
```

```
data at 'b': 0.5
```

A Series looks a little like a Dictionary, but with useful operations like NumPy ndarrays

```
[5]: data.sum()
```

```
[5]: 2.5
```

You can access the index and the values directly. The values are a NumPy ndarray so you can easily integrate with NumPy.

The Index is its own type.

```
[6]: data.index  
  
data.values  
type(data.values)
```

```
[6]: Index(['a', 'b', 'c', 'd'], dtype='object')
```

```
[6]: array([0.25, 0.5 , 0.75, 1.  ])
```

```
[6]: numpy.ndarray
```

Symbolic indexing is super convenient ! No more “parallel” arrays of values and labels.

```
[7]: data = pd.Series( np.arange(0,3), index=["2018-01-01", "2018-01-02",  
      ↪ "2018-12-31"])  
data.index
```

```
[7]: Index(['2018-01-01', '2018-01-02', '2018-12-31'], dtype='object')
```

1.2 Pandas DataFrame

A DataFrame (note the capital “F”) looks like a table or a 2-D ndarray but with *symbolic indexing* of rows and columns.

Unlike an ndarray, the columns can be different types.

1.2.1 Constructor

There are several ways to construct a DataFrame. Common methods: - list of tuples, each tuple representing a row - list of Dictionaries, each Dictionary representing a row (key/value pairs) - A Dictionary - where keys are column names - values are the Series representing a column - the index of the Series becomes *row* names of the DataFrame

```
[8]:
```

```

prices      = pd.Series( { "FB": 150, "AAPL": 156, "AMZN": 1700, "NFLX": 340,
↪ "GOOG": 1100 } )
full_names  = pd.Series( { "FB": "Facebook", "AAPL": "Apple", "AMZN": "Amazon",
↪ "NFLX": "Netflix", "GOOG": "Google"
                          } )

stocks = pd.DataFrame( { "price": prices, "name": full_names } )
stocks

```

```

[8]:      price      name
     FB      150  Facebook
     AAPL    156    Apple
     AMZN   1700    Amazon
     NFLX    340    Netflix
     GOOG   1100    Google

```

```

[9]: ticker = "AMZN"
     print( "Ticker {t}: price: {p}, full name: {fn}".format(t=ticker,
↪ "price"],
↪ "name"]
)

```

Ticker AMZN: price: 1700, full name: Amazon

Notice how we indexed symbolically into the DataFrame: `stocks.loc[ticker, "price"]`

More on indexing.

1.2.2 Display

```

[10]: stocks.head(2)

stocks.tail(3)

```

```

[10]:      price      name
     FB      150  Facebook
     AAPL    156    Apple

```

```

[10]:      price      name
     AMZN   1700    Amazon
     NFLX    340    Netflix
     GOOG   1100    Google

```

1.2.3 Data Indexing and Selection

VanderPlas

There is more than one way to index, and it gets confusing !

Personally, I just use one-way: the `.loc` indexer

```
[11]: print("Price column via .price attribute:\n", stocks.price)

# What would happen if "price" were in both the row index and the column index ?
print("\nPrice column via named column (implicitly the column):\n",
      ↪stocks["price"])

print("\nPrice column via .loc indexer:\n", stocks.loc[:, "price"])
```

Price column via `.price` attribute:

```
FB      150
AAPL    156
AMZN    1700
NFLX     340
GOOG    1100
Name: price, dtype: int64
```

Price column via named column (implicitly the column):

```
FB      150
AAPL    156
AMZN    1700
NFLX     340
GOOG    1100
Name: price, dtype: int64
```

Price column via `.loc` indexer:

```
FB      150
AAPL    156
AMZN    1700
NFLX     340
GOOG    1100
Name: price, dtype: int64
```

Common “gotcha” Note the difference in return type for the two slightly statements

```
[12]: print("Arg. is an array with a single element")
stocks.loc[:, ["price"] ]
type(stocks.loc[:, ["price"] ] )

print("Arg. is a singleton")
stocks.loc[:, "price" ]
```

```
type(stocks.loc[:, "price" ])
```

Arg. is an array with a single element

```
[12]:      price
      FB      150
      AAPL    156
      AMZN   1700
      NFLX    340
      GOOG   1100
```

```
[12]: pandas.core.frame.DataFrame
```

Arg. is a singleton

```
[12]: FB      150
      AAPL    156
      AMZN   1700
      NFLX    340
      GOOG   1100
      Name: price, dtype: int64
```

```
[12]: pandas.core.series.Series
```

Index alignment Indices are more than a convenient feature for indexing. They have semantic meaning in that they align two Series or DataFrames

```
[13]: series1 = pd.Series( { "a": 1,   "b": 2,   "d": 4})
      series2 = pd.Series( { "a": 10,  "c": 30, "d": 44 })

      series1
      series2

      series1 + series2
      series2 + series1
```

```
[13]: a      1
      b      2
      d      4
      dtype: int64
```

```
[13]: a      10
      c      30
      d      44
      dtype: int64
```

```
[13]: a    11.0
      b     NaN
      c     NaN
      d    48.0
      dtype: float64
```

```
[13]: a    11.0
      b     NaN
      c     NaN
      d    48.0
      dtype: float64
```

Slices You can use slices in Index !

Note that, unlike in Python, the upper bound is **inclusive**

Slices are particularly useful when you have DatetimeIndex (an index consisting of dates)

```
[14]: stocks.loc["AMZN": "GOOG"]
```

```
[14]:      price    name
      AMZN   1700  Amazon
      NFLX    340  Netflix
      GOOG   1100   Google
```

Adding a column

```
[15]: stocks["sector"] = pd.Series( { "FB": "Tech", "AAPL": "Tech", "AMZN": "ConsD",
    ↪ "NFLX": "ConsD" })
      stocks
```

```
[15]:      price    name sector
      FB     150  Facebook   Tech
      AAPL   156    Apple   Tech
      AMZN  1700   Amazon  ConsD
      NFLX   340  Netflix  ConsD
      GOOG  1100   Google    NaN
```

Note the missing values.

1.3 Operations on Series and DataFrames

The usual suspects, plus vector operations like NumPy.

Note that the default `axis=0`, just like NumPy

```
[16]: stocks.loc[:, "price"].mean()
```

```
[16]: 689.2
```

1.3.1 apply

There are lots of methods (hard to remember them all). Don't forget the old standby: **apply**

The **apply** method to applies your own function, either column-wise or row-wise (depending on axis chosen)

```
[17]: def my_func(s):  
        # If you don't understand what is passed (Series or DataFrame), or whether  
        → it is row or column, try this:  
        print("Type of s is {t}, shape is {sh}".format(t=type(s), sh=s.shape))  
  
        # s is a series  
        return s.mean()  
  
stocks.loc[:, ["price"] ].apply(my_func, axis=0)
```

Type of s is <class 'pandas.core.series.Series'>, shape is (5,)

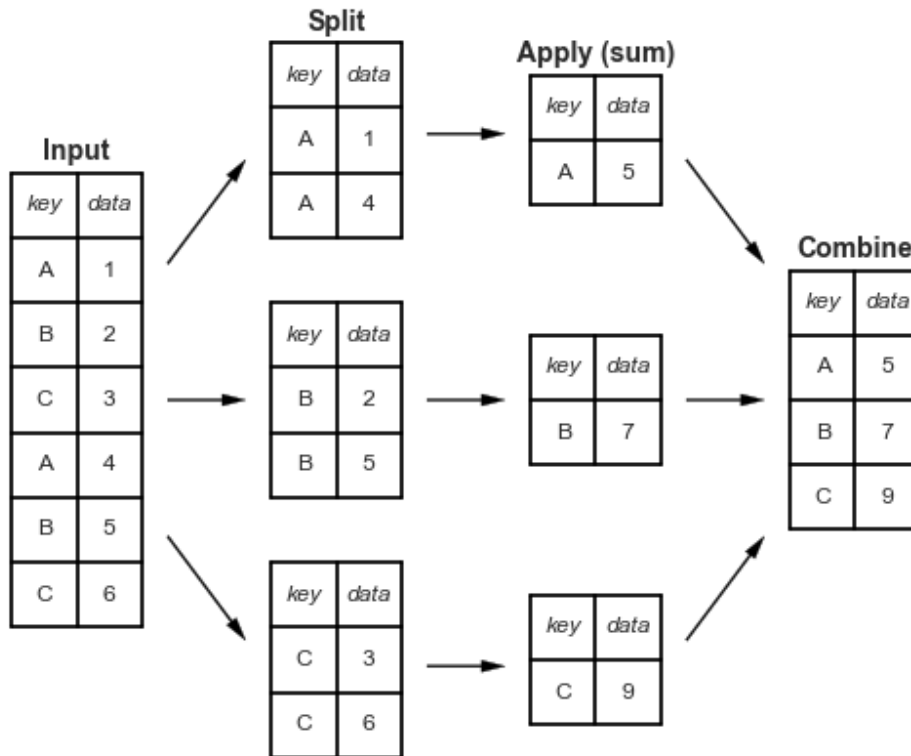
```
[17]: price    689.2  
      dtype: float64
```

1.3.2 Database (and Spreadsheet) like operations

Aggregation and grouping [VanderPlas](#)

The “split-apply-combine” paradigm ([VanderPlas](#)) should be familiar to users of SQL, where the operation is called **group by**

- split the DataFrame into groups by filtering the rows on some criteria
- apply a function to each group, returning another DataFrame
- combine the DataFrame from each group result into a single DataFrame



```
[18]: stocks.groupby("sector")
```

```
[18]: <pandas.core.groupby.groupby.DataFrameGroupBy object at 0x7f8505b34e80>
```

```
[19]: def my_group_func(df):
      return df.loc[:, "price"].mean()

stocks.loc[:, ["price", "sector"] ].groupby("sector").aggregate(my_group_func)
```

```
[19]:      price
sector
ConsD  1020.0
Tech   153.0
```

Joining, concatenation, pivot [VanderPlas](#)

It is very rare that you are given a single dataset organized exactly as you like. Often it comes in separate pieces (i.e., from different database tables or vendors) that must be joined.

You can combine two DataFrames in a manner just like an SQL join of two tables, where the join column is the Index of the two DataFrames. Remember: Index is used for alignment as well as convenience.

We will explore this in our module on Data Transformations.