

Recurrent Neural Networks

The RNN API

Sequences present several complexities.

Let's begin by better understanding functions that have sequences as inputs and output.

We call this the [RNN API \(RNN_API.ipynb\)](#).

Inside an RNN layer

By now you hopefully have a good intuitive understanding of a Recurrent Layer, but lack the details

Let's open up the hood and [go inside an RNN \(RNN Workings.ipynb\)](#).

RNN in action

A concrete example may help you to appreciate the power of an RNN.

Let's see an [RNN in action \(RNN in action.ipynb\)](#).

What is *really* going on inside an RNN

At this point

- You appreciate the ability of an RNN to operate on sequences
- Understand the mechanics of the internal workings

But the update equations don't really convey an intuition about *how* the RNN achieves its power.

Let's try to visualize the latent state of an RNN in order to get a better grasp.

[RNN Visualization \(RNN_Visualization.ipynb\)](#)

RNN practicalities

Sequences: Variable length

There are lots of small potholes one encounters with sequences.

What are the examples of my training set have widely varying lengths ?

- Within a batch, short examples may behave differently than long examples:
 - Maybe learn less in short examples, noisier gradient updates
- Padding sequences to make them equal length
 - Pad at the start ? Or at the end ?

The general advice is to arrange your data so that an epoch contains examples of similar lengths.

- You may require multiple fittings, one per length

Issues with RNN's

Although an RNN layer seems powerful (and a little magical) we have glossed over some big issues

- Can they handle *long* sequences or are they subject to "forgetting" ?
 - Short term versus long term memory tradeoffs
- Can we really unroll a computation over a long sequence ?
 - Gradient computation potentially more difficult in very deep graphs
- What are the practical difficulties in Keras with long sequences

These will be the topics of subsequent modules.

- Some topics require an in-depth understanding of Gradient Computation (still to come !)

In [2]: `print("Done")`

Done